

Scientific Programming in Julia

Performant and Elegant

Nicholas Gale & Stephen Eglen

Course Objectives

In this course we will cover:

- ▶ Why we want to learn Julia.
- ▶ How to interact with Julia and the basics.
- ▶ The key differences between Julia and R.
- ▶ Several advanced concepts: types and structures, optimisation, high performance computing.
- ▶ Some essential packages: Plotting, Statistics, Distributed, CUDA.

A new language?

- ▶ A new language needs to offer benefits over the existing language.
- ▶ Current paradigm is a two language development cycle - the *two-language problem*.
- ▶ This is not elegant and lots of code is rewritten.

Why Julia?

- ▶ Julia offers a solution to the two language problem: it is fast *and* flexible.
- ▶ It uses a method called just-in-time (JIT) compilation to generate compiled code just before it is needed.
- ▶ Easy to develop in.

Interacting with Julia

- ▶ Three primary ways to interact with Julia: Notebooks, terminal, and REPL.
- ▶ The notebook format is through the “IJulia” package and Jupyter notebook software. This is exactly like Python and R notebooks but with Julia code.
- ▶ The terminal can be used to execute scripts with the bash command `julia path/to/script`.
- ▶ REPL stands for Read-Evaluate-Print-Loop and is an interactive Julia session initiated by running `julia` in terminal.

REPL

- ▶ The REPL can execute basic commands or run scripts with the include function:

```
include("my_first_script.jl")
```

The output of my first script, is a string!

- ▶ Scripting in the REPL is costly *once*. Scripting from terminal is costly *everytime*.
- ▶ `println(arg)` is the method for printing (print works, but doesn't generate a new line).
- ▶ `display(arg)` creates a user-friendly readout in the REPL.

- ▶ Working in the REPL is very intuitive: functions generally have very mathematical names.
- ▶ Variable assignment is done through the = operator, equality is tested through the == operator, and indistinguishability through the === operator. This is different to R which uses the <- assignment operator and the identical(a,b) function for indistinguishability.

```
a = 1.0  
b = 1  
println(a == b)  
println(a === b)
```

```
true  
false
```

Operators

- ▶ Base operators are the same as in R (+, -, *, ^, etc).
- ▶ `div(a,b)` and `mod(a,b)` get division without remainder, and remainder.
- ▶ Unlike R, Julia supports infix operators: `+=`, `-=`, `*=`, `/=`.
- ▶ Infix operators modify a variable in-place and can be very useful e.g. loops:

```
a = 7  
a += 1  
a
```

8

Vectors

- ▶ A vector is created with square brackets and is column delimited as opposed to R's `c(tuple...)`:

```
a = [1, 2, 4.0, 1]
```

4-element Vector{Float64}:

1.0

2.0

4.0

1.0

- ▶ A vector can be indexed through integer, logical, and cartesian indexes using square brackets []. Ranges are covered by the : operator
- ▶ Logical indexes must be the same size as the array, unlike in R.

```
println(a[2])  
println(a[[true, false, true, false]])  
println(a[CartesianIndex(1)])  
println(a[2:4])
```

2.0

[1.0, 4.0]

1.0

[2.0, 4.0, 1.0]

- ▶ Julia doesn't support negative indexing like in R
- ▶ It does have a keyword `end` which references the end of a vector and can be combined with arithmetics to create consistent negative indexing

```
a = [1,2,3,4,5,6,7]  
println(a[2:2:(end-1)])
```

```
[2, 4, 6]
```

Attention to detail

What is this thing called a range? Beware in R:

```
x = 1:1e6  
format( object.size(x), units='MB')
```

will say the object is 3.8Mb.

whereas in Julia

```
x = 1:1e6  
sizeof(x)
```

48

This is an example of an iterator.

Matrices

- ▶ A matrix is created with semi-colons to delimit rows and spaces (or ;;) to delimit columns.

```
a = [1 2 3;  
     8 9 10;  
     4 -30 0]
```

3x3 Matrix{Int64}:

1	2	3
8	9	10
4	-30	0

- Matrices are column ordered so a linear index will access elements in the column first.

```
a = [1 2 3; 8 9 10; 4 -30 0];  
println(a)  
println(a[4])  
println(a[CartesianIndex(2,1)])  
println(a[2, 3])  
println(a[1:2, 2:3])
```

[1 2 3; 8 9 10; 4 -30 0]

2

8

10

[2 3; 9 10]

Multi Dimensional Arrays

- ▶ Generally, ; concatenates in the first dimension ; ; in the second and so on...
- ▶ Higher order arrays can be constructed like this, but it is generally cumbersome.
- ▶ Multidimensional indexes extend the rules for vectors and matrices: [dim1, dim2, dim3...]

```
a = [1; 2;; 3; 4;; 5; 6;;;
      7; 8;; 9; 10;; 11; 12]
display(a)
println(a[10])
println(a[CartesianIndex(2, 3, 2)])
```

2×3×2 Array{Int64, 3}:

[:, :, 1] =

1	3	5
2	4	6

[:, :, 2] =

7	9	11
8	10	12

10

12

General Arrays and Construction

- ▶ In Julia ranges can be created with steps like in R and are iterables but they are non-allocating (in R they create a vector).

```
a = 2:50:300  
println(a[4])  
println(a)  
println(typeof(a))
```

152

2:50:252

StepRange{Int64, Int64}

Collect

- ▶ Vectors can be created with the `collect` function which works on any iterable collection.

```
a = collect(2:50:300)
println(a)
println(typeof(a))
```

```
[2, 52, 102, 152, 202, 252]
Vector{Int64}
```

Preallocation

```
a = ones(2,3,2)
display(a)
```

2×3×2 Array{Float64, 3}:

[:, :, 1] =

1.0 1.0 1.0

1.0 1.0 1.0

[:, :, 2] =

1.0 1.0 1.0

1.0 1.0 1.0

```
b = zeros{Int64}(4,4)
display(b)
```

4×4 Matrix{Int64}:

0 0 0 0

0 0 0 0

0 0 0 0

0 0 0 0

Concatenation

- ▶ Arrays can be abstractly concatenated using space, and semicolons `;`, `;;`, `;;;`.
- ▶ Horizontal concatenation is through `hcat(array1, array2, ...)`
- ▶ Vertical concatenation is through using `vcat(arrays...)`.

```
a = ones(3,3);  
b = zeros(3,3);  
display(hcat(a,b))  
display(vcat(a,b))
```

3×6 Matrix{Float64}:

1.0	1.0	1.0	0.0	0.0	0.0
1.0	1.0	1.0	0.0	0.0	0.0
1.0	1.0	1.0	0.0	0.0	0.0

6×3 Matrix{Float64}:

1.0	1.0	1.0
1.0	1.0	1.0
1.0	1.0	1.0
0.0	0.0	0.0
0.0	0.0	0.0
0.0	0.0	0.0

Strings

- ▶ A character is indicated by apostrophes 'a'.
- ▶ A string is a partial function from indexes to character literals. Unlike R apostrophes will not work, use double quotes "a" or triple double quotes """ a """.
- ▶ Triple double quotes allow some flexibility which can be useful in writing internal scripts:

```
"""triple quotes allow "quotes" to be  
embedded.  
Multiple lines.  
"""
```

```
"triple quotes allow \"quotes\" to be\nembedded.\nMultiple lines."
```

String Indexing

- Strings can be indexed in Julia like in R and unlike R are not decoded into a character vector

```
a = "A very long string"  
println(a[2:2:10])
```

eyln

Control Flow

- ▶ True and False are given by the true and false keywords.
- ▶ Logical true and false are: 0 and 1.
- ▶ Control flow blocks always have an end to terminate them. They are initiated with a special keyword. They don't usually have curly braces like in R.

```
println(true == false)
println(true == 1)
println(false == 0)
```

```
false
true
true
```


If, else, and elsif

- ▶ If blocks can have optional elsif and else keywords inside the block:

```
cond = true
if cond
  println("If this was false, nothing would happen")
end
```

If this was false, nothing would happen

```
condelse = 1
if condelse==2
    println("It's an even prime!")
else
    println("It wasn't true. The else block executed")
end
```

It wasn't true. The else block executed

```
condif = false
condelse = "three"
if condif == 1
  println("We miss this one")
elseif condelse == "two"
  println("The first elseif block executes")
elseif condelse == "three"
  println("The second elseif block executes")
else
  println("The else block executes")
end
```

The second elseif block executes

Ternary Operator

- ▶ The ternary operator can execute if one-liners:

`conditional ? iftrue_code : else_code.`

```
res = (8 < 4) ? "Maths broke." : "Situation Normal"  
println(res)
```

Situation Normal

While

- ▶ While statements are also executed when a conditional expression evaluates as true. They can be useful in loops:

```
i = 0
while i < 10
  i += 2
  println("${i} Not done yet..")
end
println("Done")
```

```
2 Not done yet..
4 Not done yet..
6 Not done yet..
8 Not done yet..
10 Not done yet..
Done
```

For Loops

- ▶ For loops operate in the same way as R with the `for var in collection` structure.
- ▶ They can operate over any abstract collection or vector.
- ▶ The `in` operator can be replaced with `=` or `∈`.

```
for i = 1:2 println(i) end  
for i in 10:12 println(i) end  
for i ∈ 100:102 println(i) end
```

```
1  
2  
10  
11  
12  
100  
101  
102
```

List Comprehension

- ▶ Julia supports list comprehension: `[expression(i) for i in collection]` generates a vector with objects specified by expression `i`.

```
v = [i^2 for i in 1:4:13]
```

4-element Vector{Int64}:

1
25
81
169

```
a = [i*j for i in 1:3, j in 1:4]
```

3×4 Matrix{Int64}:

1	2	3	4
2	4	6	8
3	6	9	12

Functions

- ▶ Functions are similarly defined as in R, but instead of setting the function name as an object of the function method we use a function block.
- ▶ Functions are defined with the keyword `function` followed by the specification of the function name and arguments.
- ▶ Functions have a `return` keyword which will return the moment the keyword is reached.


```
function complex_modulus1(x, iy)
    tmp = x^2 + iy^2
    return sqrt(tmp)
end
```

complex_modulus1 (generic function with 1 method)

- If there is no return keyword the function will return the last line.

- ▶ Functions can be written in one line in a mathematical format.

```
complex_modulus(x, iy) = sqrt(x^2 + iy^2)
```

`complex_modulus` (generic function with 1 method)

- ▶ The recommended structure is to use a function block with a return keyword.

Splatting

- ▶ Variable arguments are given by the splat `...` operator.
- ▶ Can be used in function definition.

```
function g(x...)
  return x[end]
end
println(g(4,5,6))
```

6

- ▶ Splatting can also be used in calls to functions with multiple arguments.

```
function f(x1, x2, x3)
  return (x1*x2)-x3
end
fvec = [5,4,1.0];
println(f(fvec...))
```

19.0

- ▶ Unlike R ... must be attached to a variable name

Anonymous Functions

- ▶ Anonymous functions in Julia are given using the \rightarrow operator ('-' followed by '>').
- ▶ Functions are first class objects in Julia so these may be assigned variable names or exist as standalones

```
map(x→x^3, [2, 5, 4])
```

3-element Vector{Int64}:

8

125

64

```
F = x → x^2  
println(F(2))
```

4