

Scientific Programming in Julia

Plots and High Performance Computing

Nicholas Gale and Stephen Eglen

Plots

- ▶ Plotting is typically done through the `Plots` package.
- ▶ Plotting in Julia is relatively immature: it requires a backend.
- ▶ Select one of `PyPlot`, `Plotly`, `GR`, ... (I normally go for `GR` for speed+features).
- ▶ There is a high time-to-first plot time but it is smooth from there.

Plotting grammar

- ▶ Plotting follows a predictable grammar through the `plot` API call.
- ▶ Plots are arranged in `x data, y data` formats. Series are delineated with spaces.
- ▶ Optional keywords denote styling and can be shorted:
 - 1) `seriestype/st` = {"scatter", "line", "contour", etc.}
 - 2) `color/c` = {`:red`, `RGB(0,1,0)`, etc} (can be input as a list)
 - 3) `label` = "series label"
 - 4) `xlabel=` , `title=` , `ylabel=`


```
using Plots
```

```
pyplot()
```

```
x = collect(0:0.1:10)
```

```
y = sin.(x)
```

```
plot(x, y; st="line", label=false, title="Plot of a sin wave", y
```

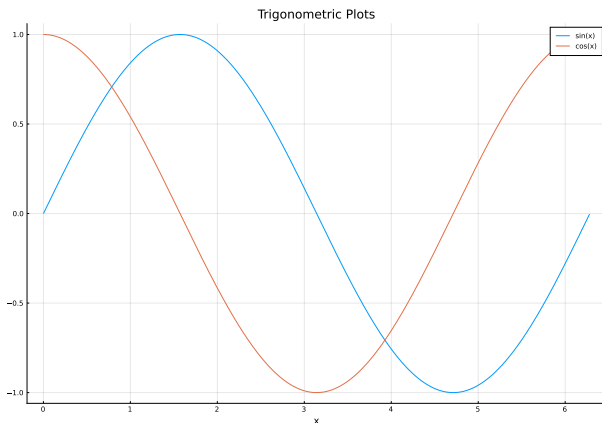


Plot Objects

- ▶ Plots can be attached to objects; last plot is attached to an internal object.
- ▶ Plot objects can be added to with the `'plot!(plot_object, x,y...; keywords=...)` call.
- ▶ Plot objects can be referenced later.

```
x = 0:0.01:2pi  
y1 = sin.(x)  
y2 = cos.(x)
```


```
trig_plot = plot(x, y1; title="Trigonometric Plots", xlabel="x"  
plot!(trig_plot, x, y2; label="cos(x)"))
```



Plot Layouts

- ▶ Subplots are handled through layout specification.
- ▶ Simple layouts can specify a grid format. A more complicated layout can be defined with the `@layout`
- ▶ A layout is arranged as a matrix with `a,b,c` denoting the plots. They are optionally given widths and heights e.g. `a{0.2h 0.3w}`.

```
x = 0:0.1:2pi  
p1 = plot(x, sin.(x)); p2 = plot(x, sin.(2x));  
p3 = plot(x, cos.(x)); p4 = plot(x, cos.(2x));  
L = @layout [a{0.2h, 0.4w} b; c d]  
plot(p1, p2, p3, p4; layout=L)
```




StatsPlots

- ▶ Plots supports recipes and thus a Plotting ecosystem.
- ▶ StatsPlots allows for efficient DataFrames plotting.
- ▶ A slightly more complex grammar; documentation found [here](#)

using DataFrames, CSV, StatsPlots


```
experiment = CSV.File("plotting_data.csv") ▷ DataFrame  
@df experiment density(:BodyLength, group = (:Planet, :Diet),  
                        legend=:topright)
```



Save Plots

- ▶ Plots are saved through the `savefig` API call.
- ▶ It specifies the path and plot object to be saved.
- ▶ The path defines the file type.


```
plt = plot(collect(1:10), rand(10); st=:scatter,  
           markersize=20)  
savefig(plt, "my_first_plot.png")
```



Animate

- ▶ A useful tool to visualise data is animation.
- ▶ This is done through the `@animate` macro applied to a `for` loop.
- ▶ The animation can be saved through the `gif` API call. This takes the animation object, then path, and optional FPS.

```
x = collect(0:0.05:2pi)
y = sin.(x)
anim_obj = @animate for t in 1:length(x)
    plot(x[1:t], y[1:t]; xrange=(0,2pi), yrange=(-1,+1),
        legend=false, xlabel="x", ylabel="sin(x)")
end
gif(anim_obj, "my_first_animation.gif", fps=30);
```



Performance Computing

- ▶ Efficient and fast code is one of the **big** draw cards of Julia.
- ▶ It *usually* comes for free, but not always.
- ▶ There are some common 'tricks' to employ to improve code efficiency and speed.

Data Access Patterns

- ▶ Data in Julia is organised in a *column-major* format.
- ▶ Data is laid out with the columns stacked end-to-end. The fastest way to access them is through the *rows*.
- ▶ Slow code uses columns as the 'fast-changing' index, but is more traditional. When looping use columns as the outer variable.


```

using BenchmarkTools

l = 10^4
A = rand(l,l);

sum1 = function(A, l)
    sum = 0.0;
    for i = 1:l, j=1:l
        sum = sum + A[i,j]
    end
    sum
end

s = sum(A); s1 = sum1(A, l)
(s ≈ s1) || error("s1 wrong")
@btime sum1(A, l);

```

237.033 ms (1 allocation: 16 bytes)

```

# .
# .
# .
# .
# .

sum2 = function(A, l)
    sum = 0.0;
    for j = 1:l, i=1:l
        sum = sum + A[i,j]
    end
    sum
end

s2 = sum2(A, l)
(s ≈ s2) || error("s2 wrong")
@btime sum2(A, l);

```

93.042 ms (1 allocation: 16 bytes)

@inbounds

- ▶ Bounds checking is a useful operation in an interpreted language throwing an error if accessing an invalid index.
- ▶ It is expensive and fast languages like C don't use it - at user peril!
- ▶ It is activated with the `@inbounds` macro and distributes to all nested loops.
- ▶ Can also use command-line switch `'-check-bounds={yes|no|auto*}'`.

@SIMD

- ▶ SIMD stands for single instruction multiple data and is a machine level optimisation in modern CPUs.
- ▶ It allows many mathematical operations to be vectorised and optimised within the CPU cycle.
- ▶ Turn it on using the '@simd macro. For well organised data you can expect some speed up.
- ▶ See extended example `timeit.jl` in repository. On my Macbook pro, it went from 2.3 GFlop/s to 18.1 GFlop/s with @inbounds and @simd in tandem.

Type mutation

- ▶ Julia has excellent typing and type inference but abstraction puts demands on the compiler.
- ▶ When types mutate the compiler works hard to “deal with it” which is good for the end user.
- ▶ To avoid this Julia can strictly type variables and functions which alleviates the pressure.

```
a::Array{Float32, 1} = [4.0, 2.0, 1.2] # a strictly typed vector  
f(x::Array{Float32, 1}) = sum(x) # a strictly typed function.
```

```
f (generic function with 1 method)
```

High Performance Computing

- ▶ HPC refers to distributing compute tasks in an efficient way.
- ▶ It typically refers to parallel computing which comes in two general flavours: multiple CPU threads, or GPUs.
- ▶ These follow similar principles but the architectures require different coding styles.
- ▶ Julia abstracts many of the 'gotchas'.

Distributed

- ▶ Parallel computing with multiple threads is available through the Distributed package.
- ▶ It supports low level (spawn, fetch, remotecall etc.) methods.
- ▶ More often it is used to parallelise loops through the @distributed macro before a for loop block.
- ▶ To make processes available use the addprocs(n) method.

```
using Distributed
addprocs(4)

@time for i = 1:10^9
    i^2
end

@time @sync @distributed for i = 1:10^9
    i^2
end
```

1.999521 seconds

0.974476 seconds (72.52 k allocations: 3.535 MiB, 2.38% compilation time)

Task (done) @0x0000000147a46290

Distributed Reduce

- ▶ A useful operation for parallelisation is reduction.
- ▶ A distributed for block can specify a binary operation to reduce on.

```
sumodds = @distributed (+) for i = 1:100  
    Int(isodd(i) && i)  
end
```

2500

Shared Functions

- ▶ Each thread has access to its own local environment and thus function definitions.
- ▶ The `@everywhere` macro is used to specify that a function/package can be accessed from thread.

```
using Distributed
@everywhere function myfunc(x)
    return x^2 - y^3 + sin(x * y)
end

@everywhere using Statistics
```

Shared Data

- ▶ Additional threads also do not have access to data on the master thread.
- ▶ The SharedArrays package allows for memory to be shared between threads through a SharedArray object.

```
using Distributed
@everywhere using SharedArrays

v = SharedArray(zeros(5))

@sync @distributed for i = 1:length(v)
    v[i] = i^2
end
println(v)
```

```
[1.0, 4.0, 9.0, 16.0, 25.0]
```

Race Conditions and Synchronisation

- ▶ Data and thread access is typically asynchronous.
- ▶ What happens when two threads depend on each other catastrophically?
- ▶ What happens when two threads try to access/modify data at the same time.
- ▶ These are *race* conditions and the second can be dealt with using *atomic* operations and synchronisation.

pmap

- ▶ Functional programming styles are supported through pmap which behaves like map.
- ▶ There is a slight performance cost and is best used with complex function calls.

```
pmap( x → x^2, 1:5)
```

5-element Vector{Int64}:

1
4
9
16
25

GPU Compute

- ▶ Graphical Processing Units use a special hardware layout to launch hundreds of threads with a low clock speed.
- ▶ They are slow at performing individual tasks, but can do many simultaneously.
- ▶ The startup time is slow but for large jobs speed ups of 20x to 100x are commonplace.
- ▶ Julia supports them through a generic GPUArrays backend, but users interface with a card specific API: AMDGPUs, CUDA, Metal.
- ▶ They are best used for linear algebra and vectorisable tasks.

GPU Basics

- ▶ APIs have a shared array data structure which can load CPUs to the GPU: `CUDA.CuArray`, `AMDGPU.ROCArray`, `Metal.MtlArray`.
- ▶ Common functions are overloaded to these data structures: `sum`, `+`, `-`, `*`, `/`, `^`, `..`, `sin`, `cos`, `exp` etc.
- ▶ Functional programming is supported through map and reduce frameworks.
- ▶ Indexed functions are *highly* discouraged: the GPU has to fall back on the CPU incurring a *very* high cost.
- ▶ Calls are asynchronous and the `PACKAGE.@sync` macro will synchronise data which is important for data dependencies.

```
using CUDA
a = rand(10^8);
a_gpu = CuArray(a);
@time sum(a);
CUDA.@time CUDA.@sync sum(a_gpu)
```

Further reading

- ▶ Performance tips:
<https://docs.julialang.org/en/v1/manual/performance-tips/>
- ▶ Static arrays can often be faster (for small arrays up to about 100 elements):
<https://github.com/JuliaArrays/StaticArrays.jl>
- ▶ Floops.jl for more advanced speed ups – see Alan Edelman's birthday problem video:
<https://www.youtube.com/watch?v=dczkYIOM2sg>
- ▶ Julia style guide:
<https://docs.julialang.org/en/v1/manual/style-guide/>

Summary

- ▶ Plotting
- ▶ Performance
- ▶ High-performance computing